

Bloom filters

Introduction

How to store a set, defined with following methods:

```
set.add(x)
set.remove(x)
set.contains(x)
```

Ideas:

1. Bit set (every bit for a specific property), efficient, can only store pre defined fixed numbers
2. Hash table (hash sets)
Tree (balanced)
 - Fast $O(1)$
 - Require a lot of memory, objects are complex
3. Trie (Prefix trees, for strings)
4. Bloom filters
 - if $x \in S$, return *yes*
 - if $x \notin S$, return with $P_{no} \rightarrow 1, P_{yes} \rightarrow 0$
 - Do not support removing elements

Naive bloom filters

(Assumption: set is sparse, compared to element universe u .)

We want to store m elements, with k bits per element.

Let size $n = mk$ (total size in bits)

T : Bit set of size n (the underlying storage used to represent values)

h : hash function (from universal hash family / ROM), where $h : u \rightarrow \{0, \dots, n\}$

Define the following methods for set:

```
def add(x):
    id = h(x)
    T[id] = 1

def test(x):
    id = h(x)
    return T[id]
```

Performance definition

lemma:

Let x_1, \dots, x_m be distinct elements inserted in the naive bloom filter, y is an element not in $\{x_1, \dots, x_m\}$
then:

$$P\{test(y) = 1\} \leq 1/k$$

Proof

proof:

$$test(y) = 1 \Leftrightarrow \exists i, h(x_i) = h(y)$$

then:

$$\begin{aligned} P\{test(y) = 1\} &= P\left(\bigcup_i h(x_i) = h(y)\right) \\ &= \sum_{i=1}^m P\{h(x_i) = h(y)\} \\ &\leq \frac{1}{n} * m \\ &= \frac{1}{m * k} * m \\ &= \frac{1}{k} \end{aligned}$$

Bloom filters

h_1, \dots, h_k : hash functions (ROM, independent from each other), where

$h_1, \dots, h_k : u \rightarrow \{0, \dots, n\}$

Let size $n = mk$

Define the following methods for set:

```
def add(x):
    for i = 1...k:
        T[h_i(x)] = 1

def test(x):
    return all(T[h_i(x)] = 1 for all i)
```

Performance definition

lemma:

Let x_1, \dots, x_m be distinct elements inserted in the naive bloom filter, y is an element not in $\{x_1, \dots, x_m\}$

then:

$$P\{test(y) = 1\} \leq (1 - 1/e)^k$$

Informal proof

Proof:

$$P\{test(y) = 1\} = P\{\forall i, T[h_i(y)] = 1\}$$

The RHS of this equation could be decomposed into:

$$\begin{aligned}
& P\{\forall i, T[h_i(y)] = 1\} \\
&= \prod_{i=1}^k P\{T[h_i(y)] = 1\} \\
&= (P\{T[h(y)] = 1\})^k
\end{aligned}$$

Note: The independence assumption of this transformation is **incorrect**, we will cover this part in the next section

Therefore, the probability of one of k hash values of y having no collision in table T is:

$$\begin{aligned}
& P\{T[h(y)] = 0\} \\
&= P\{\forall i, j (i = 1 \dots k, j = 1 \dots m), \quad h_i(y) \neq h_i(x_j)\} \\
&= \left(1 - \frac{1}{n}\right)^{mk} \\
&= \left(1 - \frac{1}{mk}\right)^{mk} \\
&= 1/e \quad \left(\text{when } \lim_{mk \rightarrow \infty}\right)
\end{aligned}$$

With this result we can compute the value of probability $P\{\forall i, T[h_i(y)] = 1\}$:

$$\begin{aligned}
& P\{\forall i, T[h_i(y)] = 1\} \\
&= \prod_{i=1}^k P\{T[h_i(y)] = 1\} \\
&= (P\{T[h(y)] = 1\})^k
\end{aligned}$$

Note: $O(1)$ is a remainder, since mk is not ∞

$$= \left(1 - \frac{1}{e} + O(1)\right)^k$$

Therefore the upper bound of the probability of a false positive (where all $T[h_i(y)] = 1$) is:

$$P\{\text{false positive}\} \leq (1 - 1/e)^k$$

Why informal proof is incorrect

More detailed analysis could be found in this paper: [On the false-positive rate of bloom filters](#)

Lets enumerate all positive conditions in a simple example, where $m = 1, k = 2, n = mk = 2$

Let A be the hash value of h_1 , B be the hash value of h_2 , we have the following combination table:

bits in T		bits of h(y)	
B	A	B	A
B	A	A	B
B	A	AB	-
B	A	-	AB
A	B	B	A
A	B	A	B
A	B	AB	-
A	B	-	AB
AB	-	B	A
AB	-	A	B
AB	-	AB	-
AB	-	-	AB
-	AB	B	A
-	AB	A	B
-	AB	AB	-
-	AB	-	AB

As we can see:

$$P\{h_1(y) = 1\} = P\{h_2(y) = 1\} = 12/16 = 3/4$$

$$P\{h_1(y) = 1, h_2(y) = 1\} = 10/16 = 5/8$$

These two probabilities are not independent in this case, because when table T is fully occupied, $P\{h_1(y) = 1, h_2(y) = 1\}$ becomes larger, in the upper 8 rows, this probability is 1, while in the lower eight rows, this probability is 1/4, which is exactly $1/2 * 1/2$, and $P\{h_1(y) = 1\} = P\{h_2(y) = 1\} = 1/2$ in the lower eight rows, so independence only stands when the load of table T is not close to 1.

However, please take note that **dependence** only exists between $P\{h_i(y) = 1\}$ because we are **not informed** about how bit table T is occupied. When T is given, probabilities $P\{h_i(y) = 1|T\}$ are **independent**, because $h_i(y)$ are independent from each other, and they are also independent from T . We are going to use this independence to construct the correct, formal proof.

Formal proof

Let $n = 2mk$ in this case (In fact, 2 could be any number larger than $\frac{e}{e-1} \simeq 1.582$)

Theorem:

After inserting x_1, \dots, x_m in Bloom filter, the upper bound of false positive probability is constrained by:

$$P\{test(y) = 1\} \leq \frac{1}{2^k} = \frac{1}{2^{n/2m}} < (1 - 1/e)^{n/2m}$$

Proof:

Let \mathcal{S} be the set of bits set after we insert x_1, \dots, x_m , there is inequality:

$$|\mathcal{S}| \leq mk = n/2$$

For false positive probability there is:

$$\begin{aligned} & P\{test(y) = 1\} \\ &= P\{\forall i, T[h_i(y)] = 1\} \\ &= P\{h_1(y), h_2(y), \dots, h_k(y) \in \mathcal{S}\} \\ &= \mathbb{E}_{\mathcal{S}}[P\{h_1(y), \dots, h_k(y) \in \mathcal{S} | \mathcal{S} = S\}] \quad (\text{expectation of probability when given } S) \end{aligned}$$

if the second term below, i.e. the conditional probability is proven $\leq 1/2^k$

$$= \sum_{S'} P\{S = S'\} * P\{h_1(y), \dots, h_k(y) \in \mathcal{S} | \mathcal{S} = S\}$$

then:

$$\leq 1/2^k$$

Now goal becomes proving: $P\{h_1(y), \dots, h_k(y) \in \mathcal{S} | \mathcal{S} = S\} \leq 1/2^k$

$$P\{h_i(y) \in \mathcal{S}, \forall i | \mathcal{S} = S\}$$

(Note: These probabilities are independent, as mentioned above)

$$\begin{aligned} &= \prod_{i=1}^k P\{h_i(y) \in \mathcal{S} | \mathcal{S}\} \\ &= \prod_{i=1}^k \frac{|\mathcal{S}|}{n} \\ &\leq 1/2^k \end{aligned}$$

This concludes our proof.